Digit recognition

Final Presentation

Presented by Daria Morkis, Alex Kohlmann and Karolina Walach

Supervisor: Dr. Leonid Kostrykin, PD Dr. Karl Rohr

Tutor: Hannah Winter

What does our code include?

1. PCA
2. Self-implemented KNN
3. KDTree and KNeighborsClassifier
4. confusion matrix and classification report
5. SVM
6. CNN

Used Libraries

KD Tree

Fast way to calculate accuracy

→narrows down the area where the nearest neighbor is searched

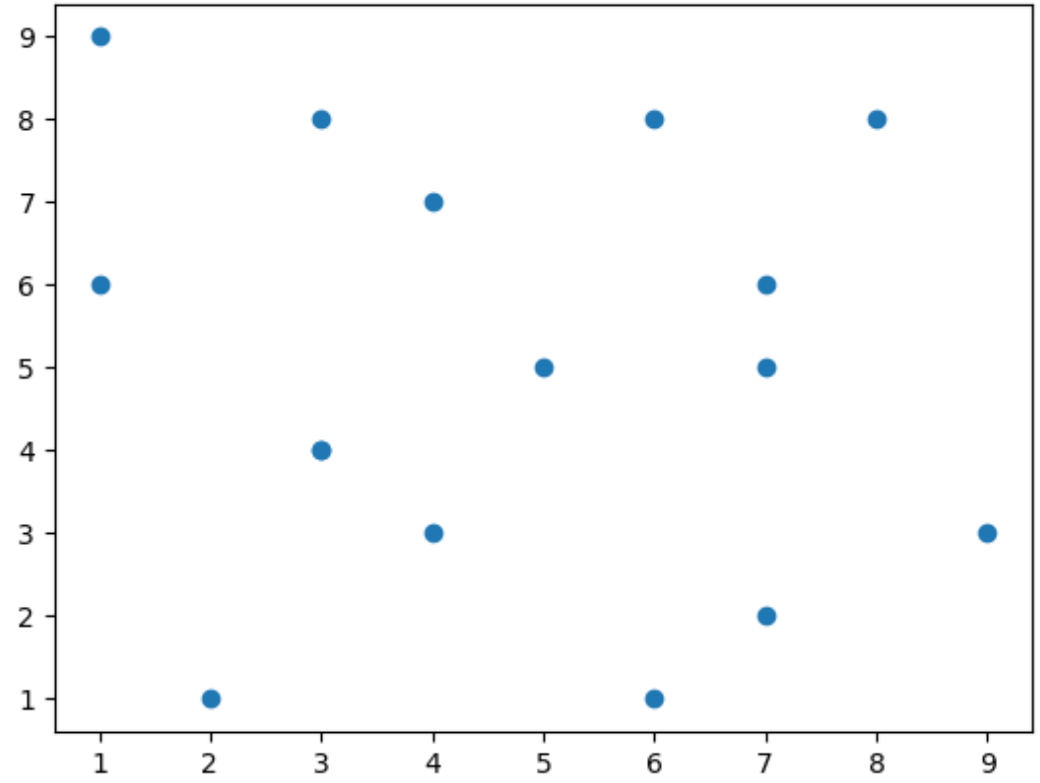

KD Tree

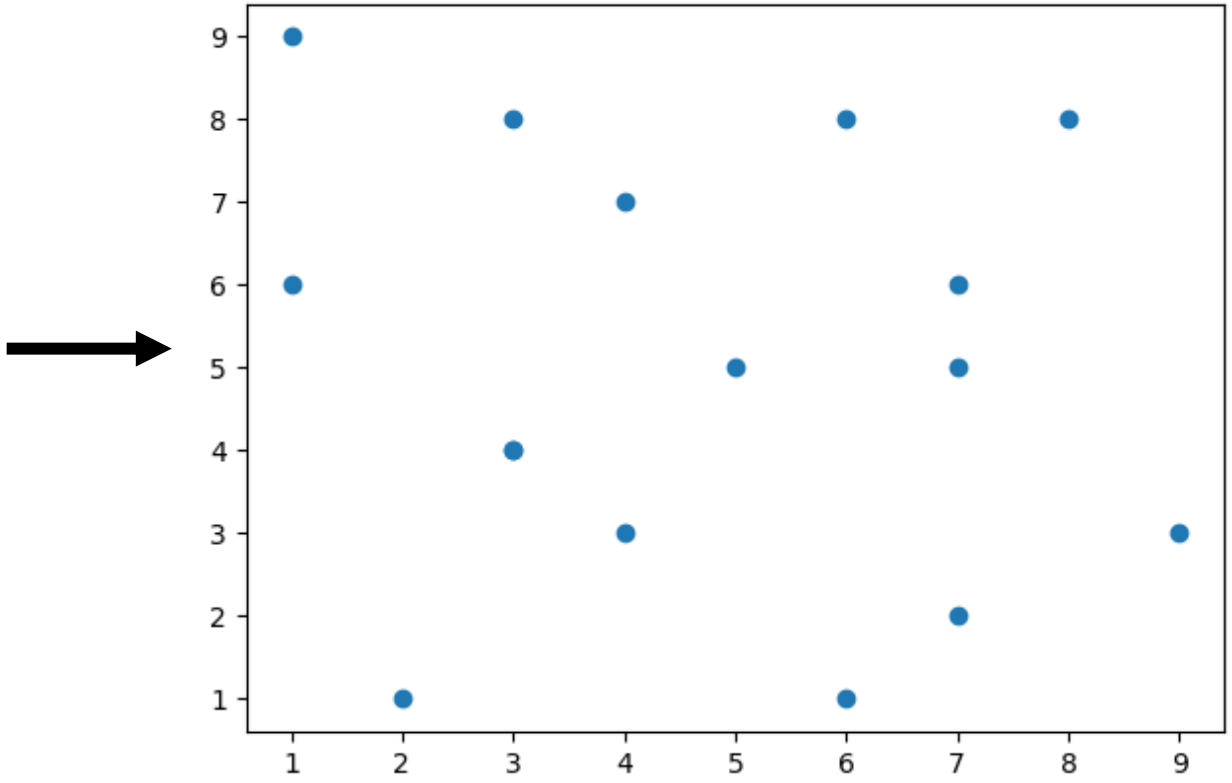

KD Tree

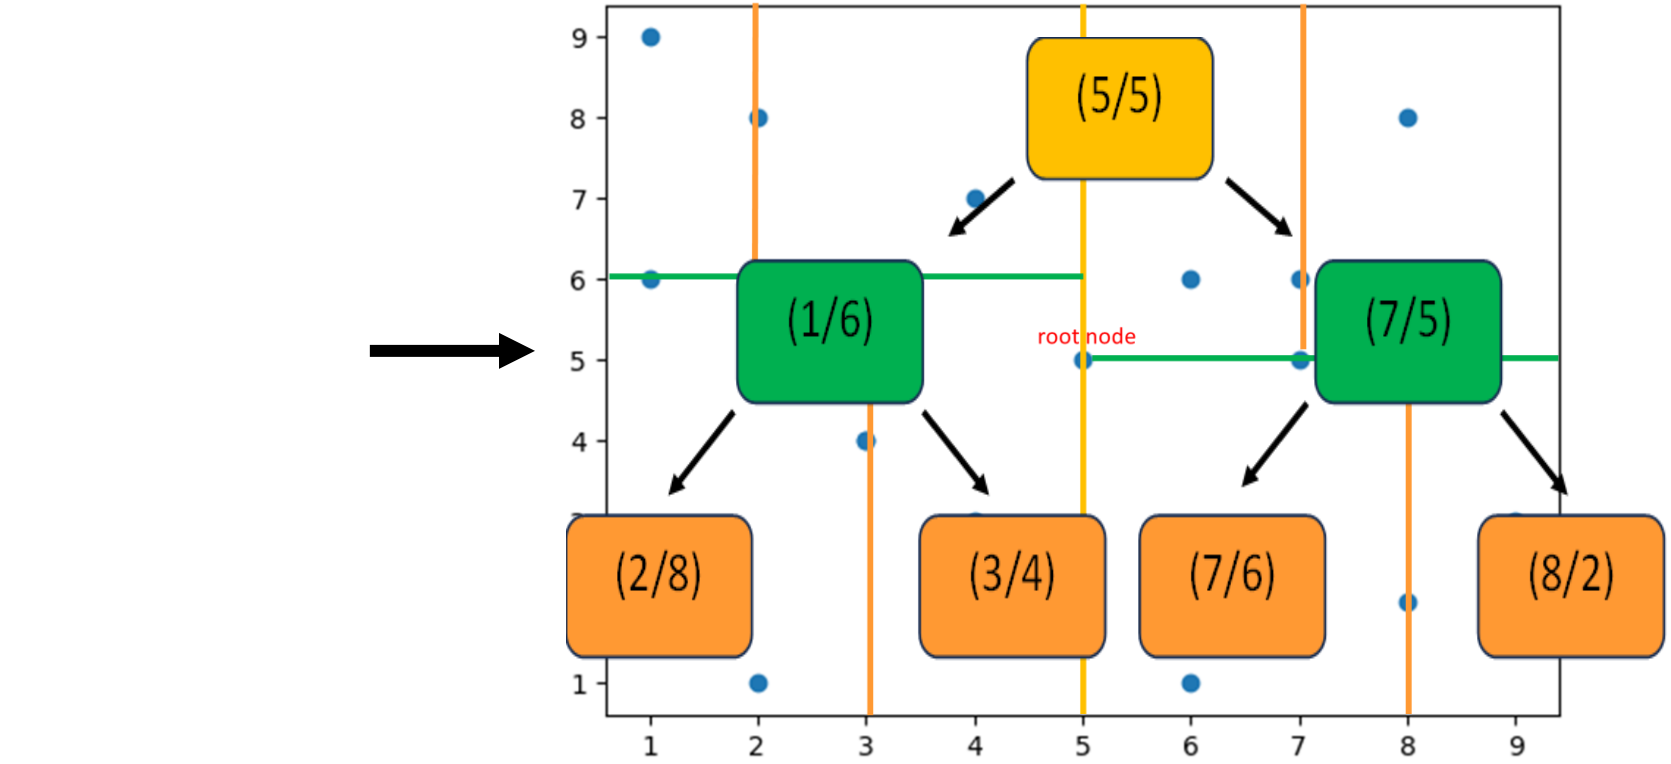

KNeighborsClassifier

```
sklearn.neighbors.KNeighborsClassifier(n_neighbors=5, weights='uniform', algorithm='auto', leaf_size=30, p=2, metric='minkowski', metric_params=None, n_jobs=None)
```

Runtime of algorithms

algorithm	runtime for **Apple M1 Max processor**		runtime for **intel core i5 processor**	
K-nearest-neighbour (self-implemented)	2.95 s	Ø 2.97 s	5.87 s	Ø 5.73 s
	2.98 s		5.72 s	
	2.97 s		5.71 s	
	2.97 s		5.67 s	
	2.97 s		5.68 s	
KD Tree (from SciPy)	3.58 ms		7.59 ms	
KNeighborsClassifier (from scikit-learn)	0.2 ms		0.3 ms	

KNN vs. KD-Tree vs. KNeighborsClassifier

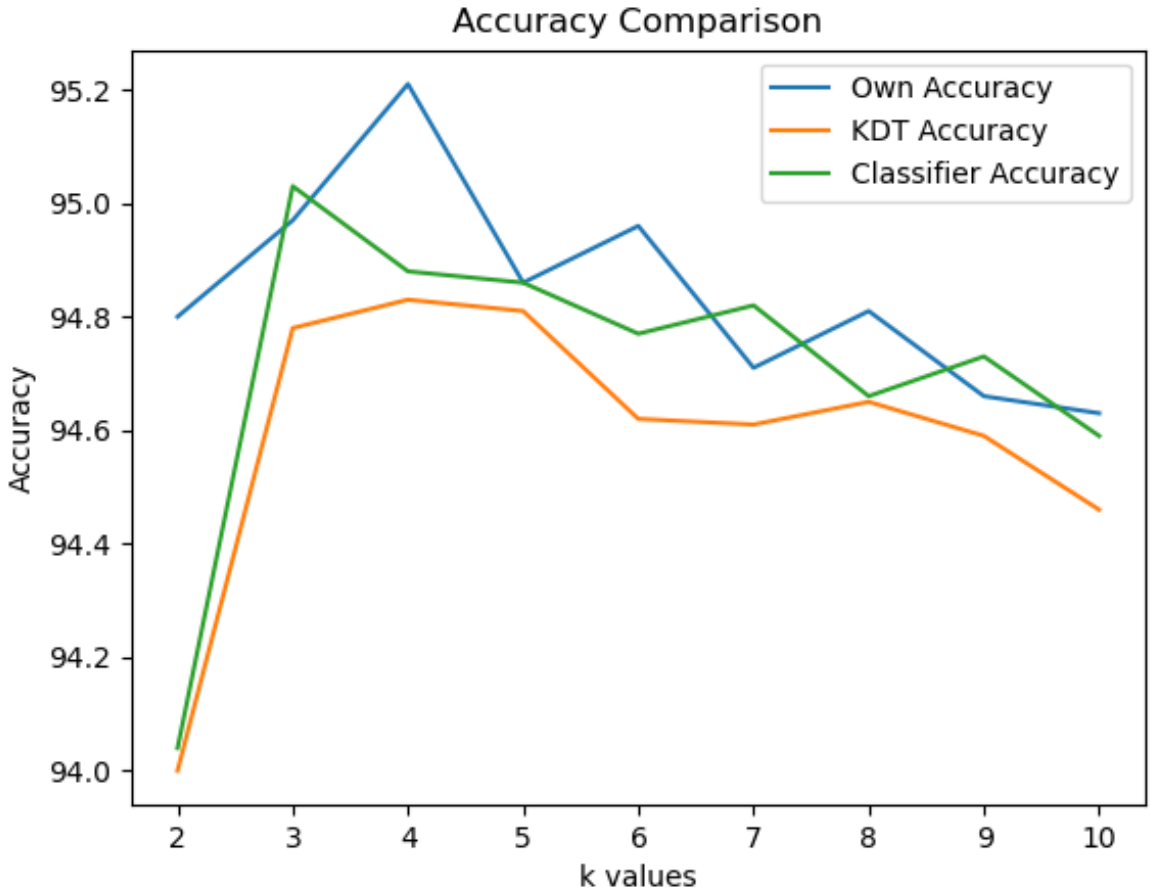

Best k-value at pc=330

KNN: k=4 with 95.21%

KD-Tree: k=4 with 94.83%

Classifier: k=3 with 95.03%

Why is k relatively small?
→ Differences of the first euclidean distances are small in ascending order

Why are the accuracies different?

Possible explanations:

1. Different rounding of euclidean distances
2. Different selection of nearest neighbor
3. Different selection of most common label (when k even)

Principal components

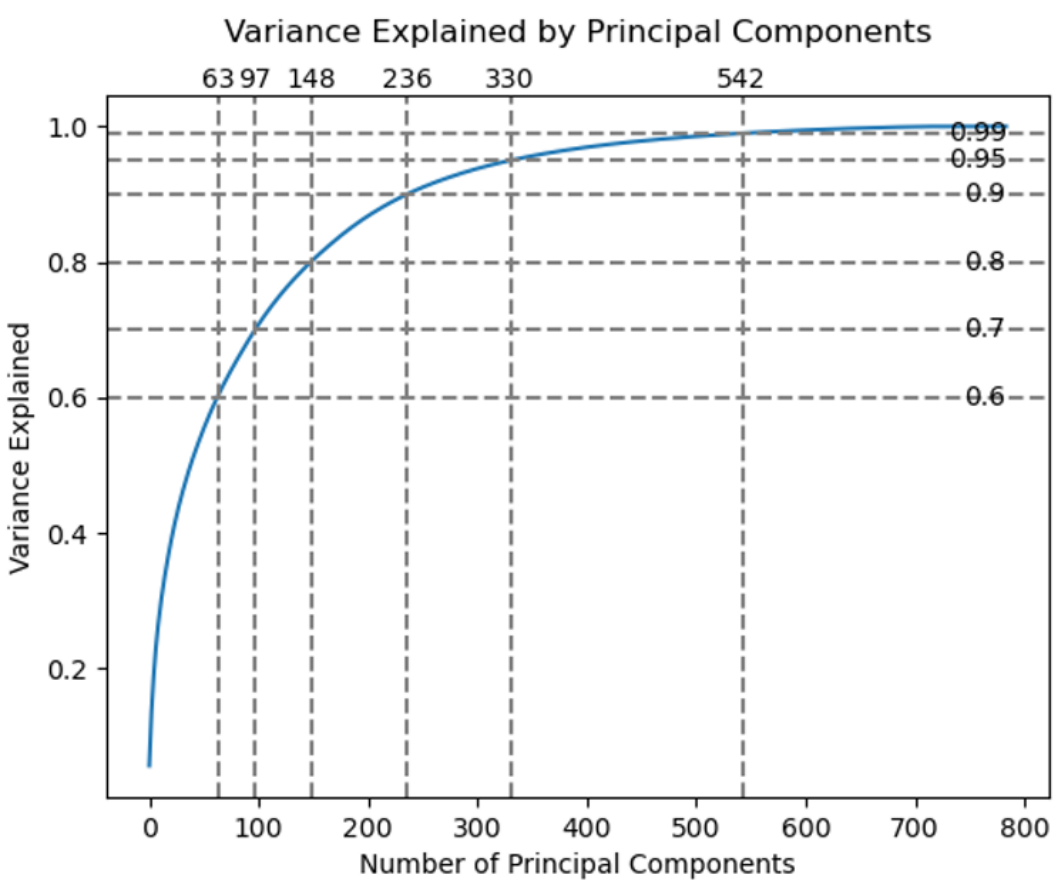

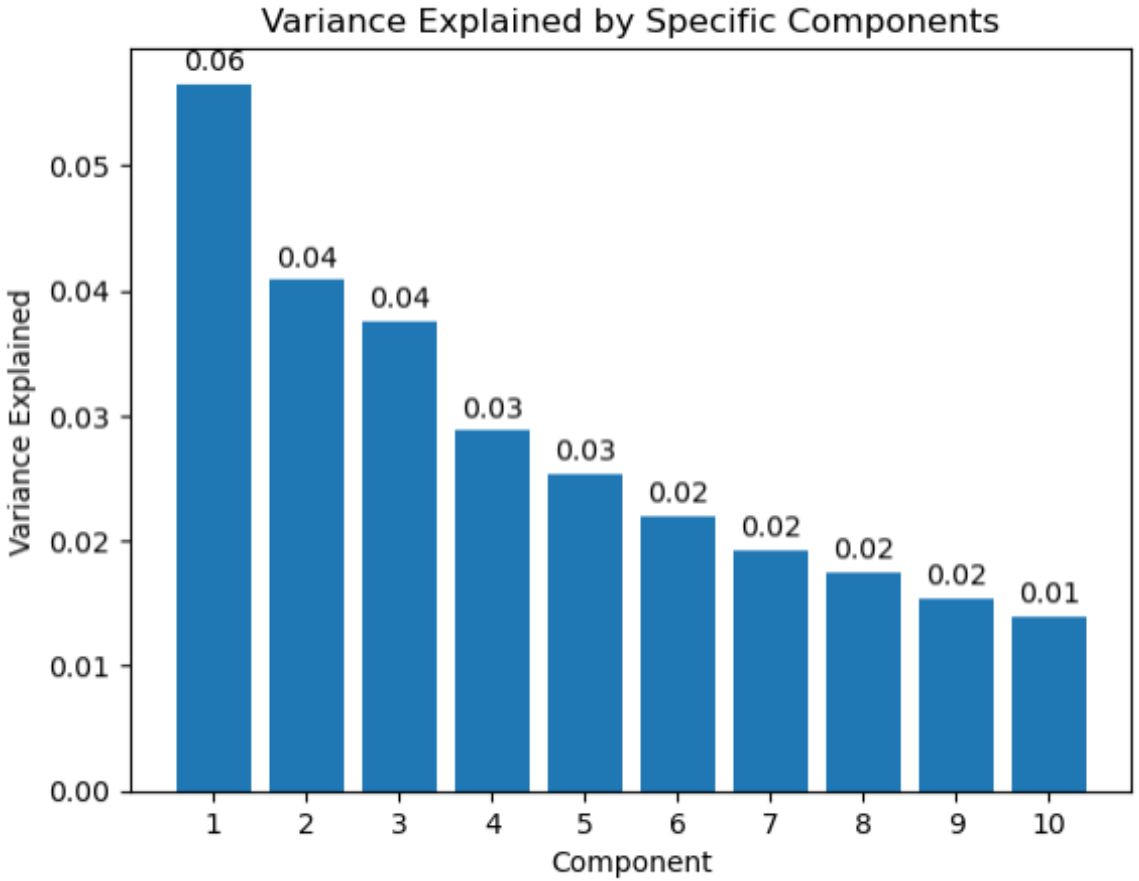

Execution time and accuracies for different numbers of PCs

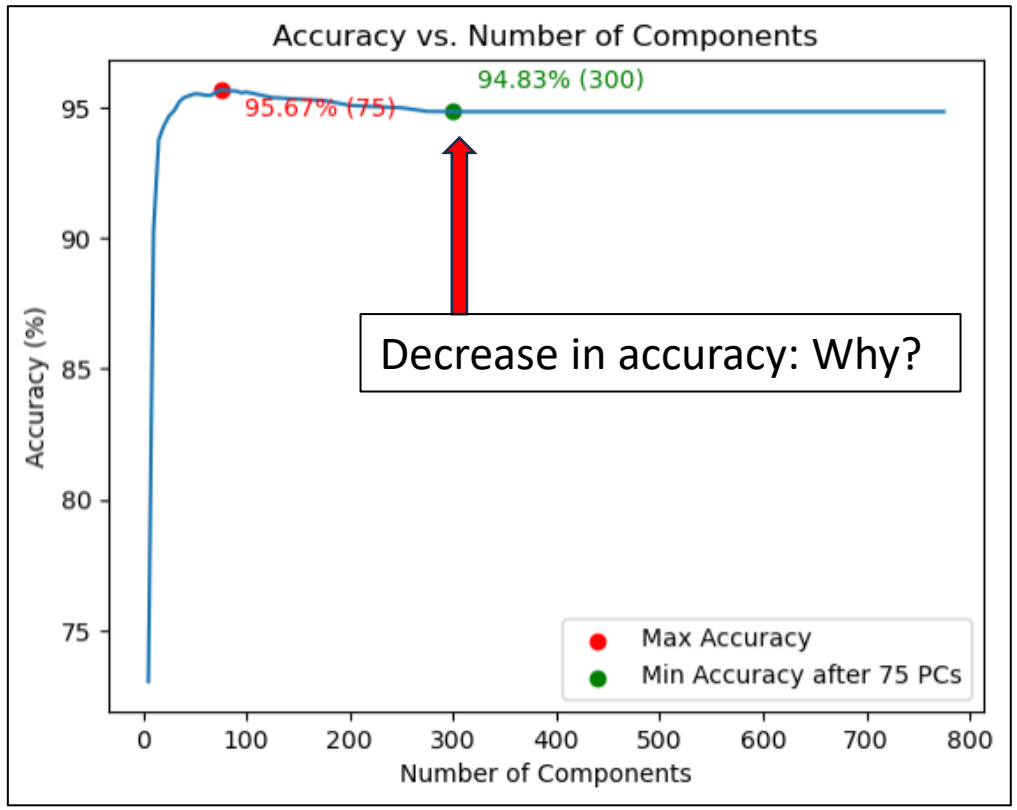

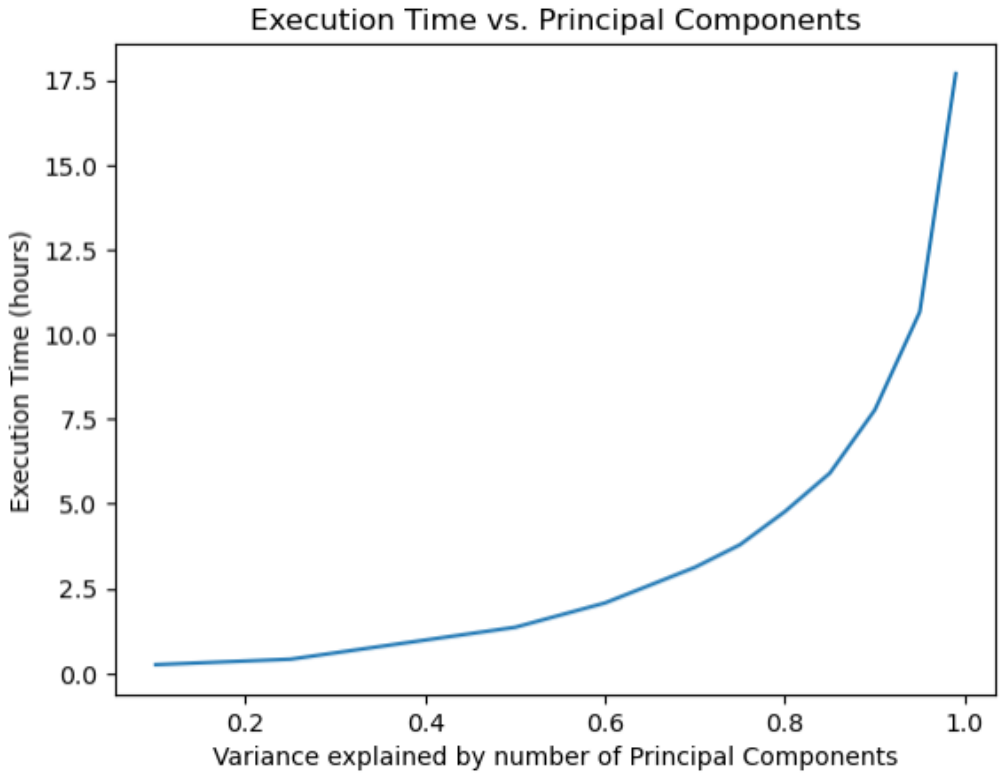

Curse of dimensionality

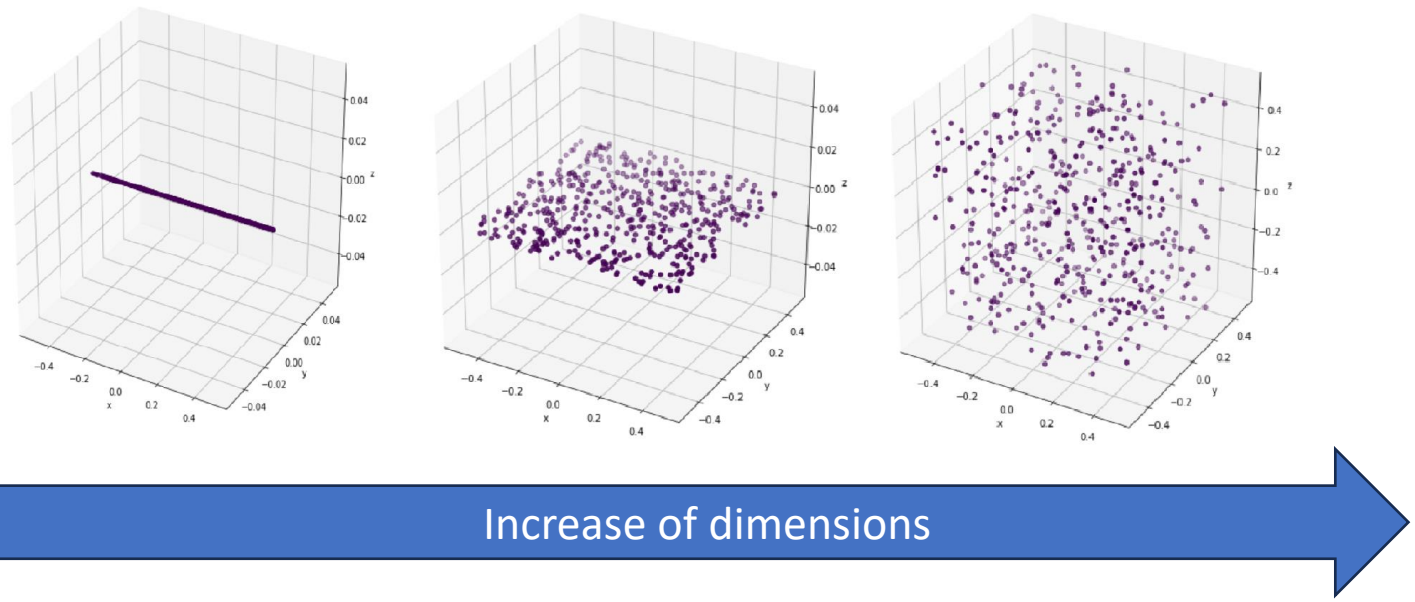

- More PCs -> more information
- More possibilities to differ from one another
- increase of distance between datapoints
- Finding k-nearest neighbors becomes more difficult

Error Analysis – balanced or imbalanced?

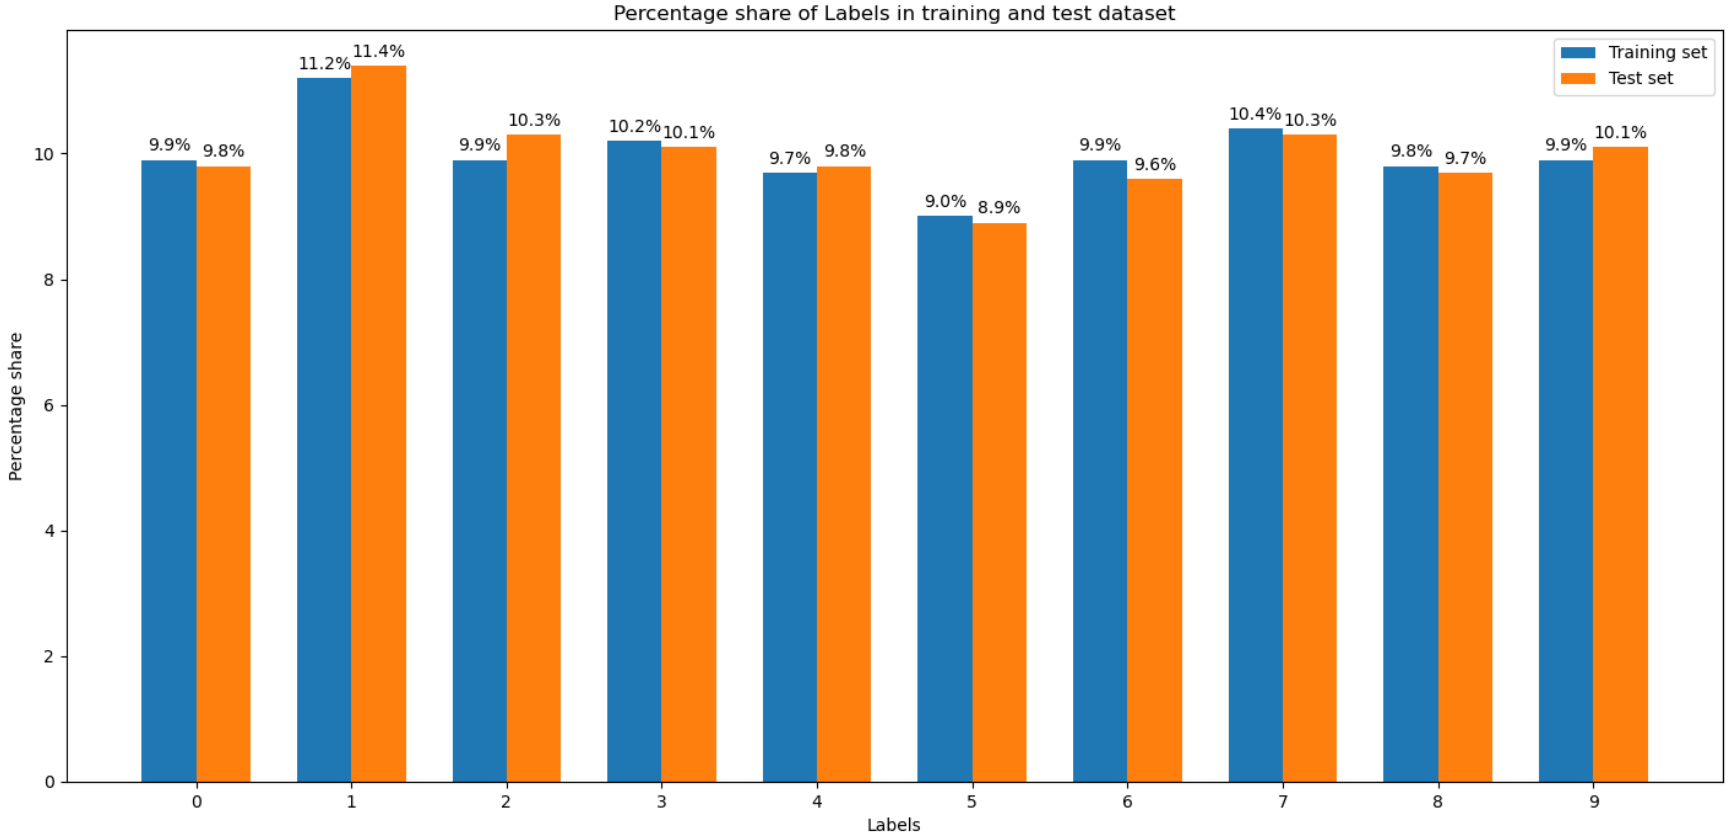

Error Analysis – classification report

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN} \qquad Precision = \frac{TP}{TP + FP}$$

$$Recall = \frac{TP}{TP + FN} \qquad F1 = 2 \times \frac{Precision \times Recall}{Precision + Recall}$$

Error Analysis – classification report

digit	precision	recall	f1-score	support
0	0.96	0.99	0.98	980
1	0.97	0.99	0.99	1135
2	0.96	0.96	0.96	1032
3	0.94	0.96	0.95	1010
4	0.96	0.96	0.96	982
5	0.94	0.95	0.95	892
6	0.97	0.97	0.97	958
7	0.95	0.94	0.95	1028
8	0.95	0.94	0.94	974
9	0.95	0.93	0.93	1009
accuracy			0.96	10000
macro avg	0.96	0.96	0.96	10000
weighted avg	0.96	0.96	0.96	10000

Error Analysis – confusion matrix

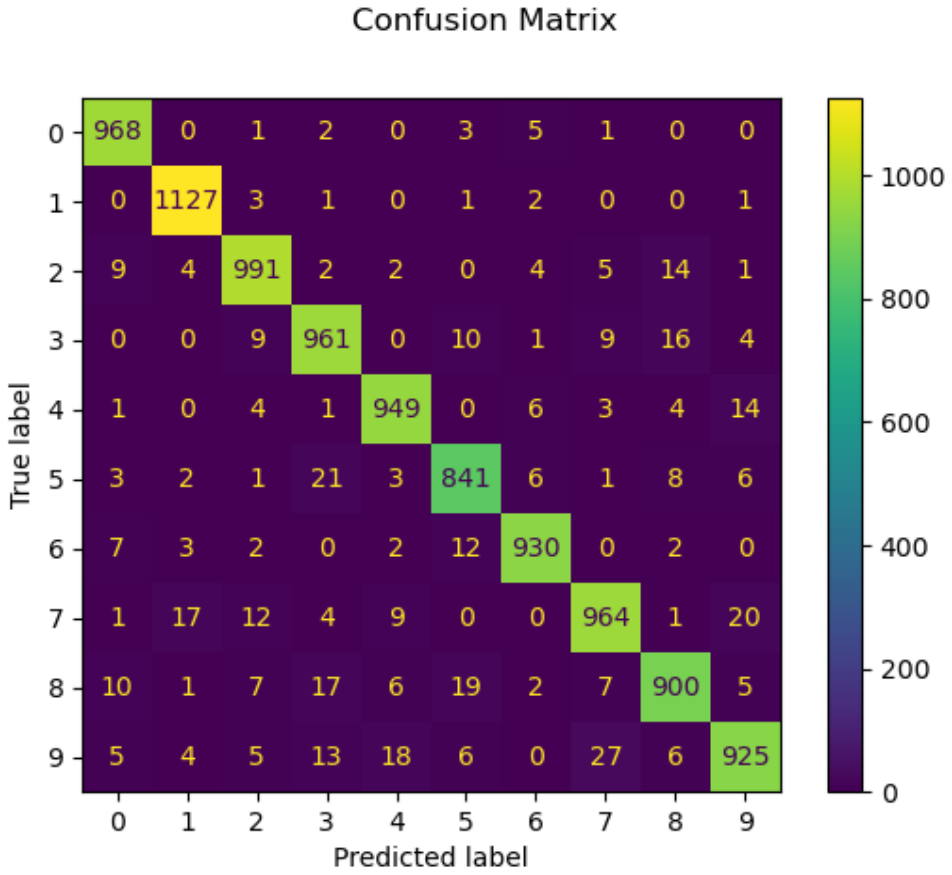

For k = 4 and variance = 0.64
Accuracy = 95.89%

Improvements – SVM

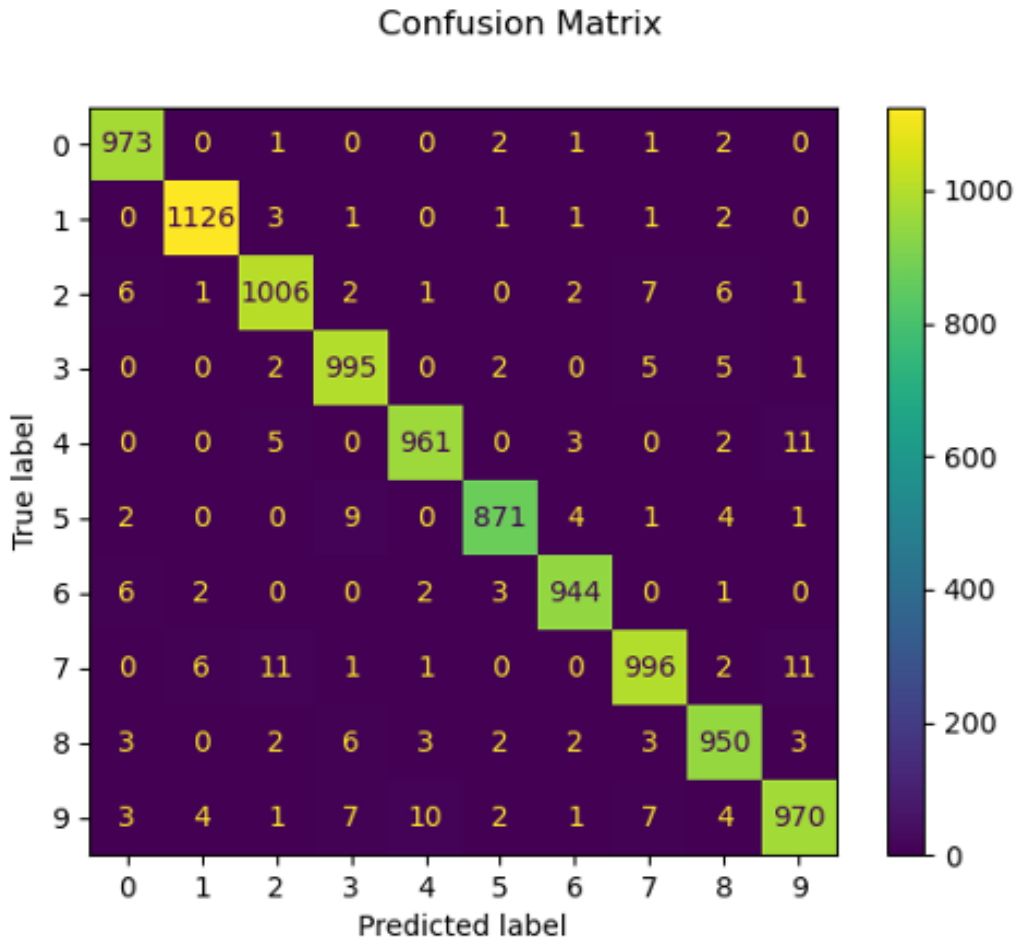

Accuracy = 97.92%

Improvements – CNN

- State-of-the-art method
- Accuracies of up to 99.80% possible
- Our CNN: 99.08%
- Reason: hierarchical feature extraction and end-to-end optimization

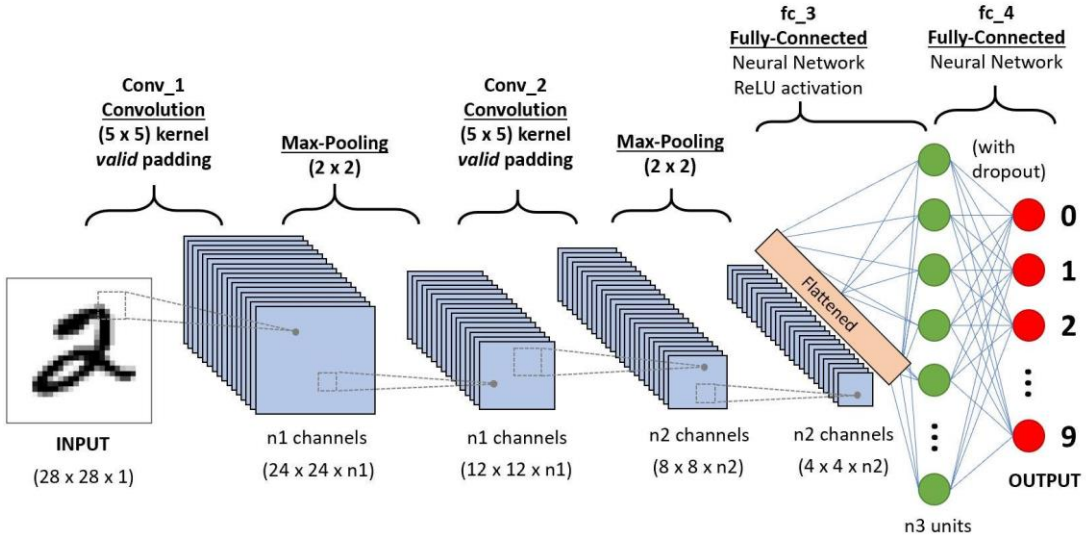

References

Grant, P. (2019). k-Nearest Neighbors and the Curse of Dimensionality. https://towardsdatascience.com/k-nearest-neighbors-and-the-curse-of-dimensionality-e39d10a6105d. accessed on: 11.07.2023

Haran, B. (2022). K-d Trees - Computerphile. https://www.youtube.com/watch?v=BK5x7IUTIyU. accessed on: 11.07.2023

Hucker, M. (2020). Tree algorithms explained: Ball Tree Algorithm vs. KD Tree vs. Brute Force.

Meigarom (2017). Dimensionality Reduction — Does PCA really improve classification outcome? https://towardsdatascience.com/dimensionality-reduction-does-pca-really-improve-classification-outcome-6e9ba21f0a32. accessed on: 12.07.2023

Géron, A. (2022). Hands-on machine learning with Scikit-Learn, Keras, and TensorFlow. O'Reilly Media, Inc.

Kanstrén, T. (2020). A Look at Precision, Recall, and F1-Score. https://towardsdatascience.com/a-look-at-precision-recall-and-f1-score-36b5fd0dd3ec. accessed on: 11.07.2023

Kasperek, D., Podpora, M., and Kawala-Sterniuk, A. (2022). Comparison of the Usability of Apple M1 Processors for Various Machine Learning Tasks. Sensors 22, 8005.

Klein, B. NumPy Tutorial. https://www.python-kurs.eu/numpy.php. accessed on: 07.07.2023

LeCun, Y., Bengio, Y., and Hinton, G. (2015). Deep learning. Nature 521, 436-444. 10.1038/nature14539.

scikit-learn. sklearn.neighbors.KNeighborsClassifier. https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html. accessed on: 13.07.2023

Additional Slides

K-Nearest Neighbors

1. Calculating euclidean distance between test data point and train data points
2. Sort distances in ascending order
3. Select top k-rows
4. Majority vote
5. Calculate accuracy

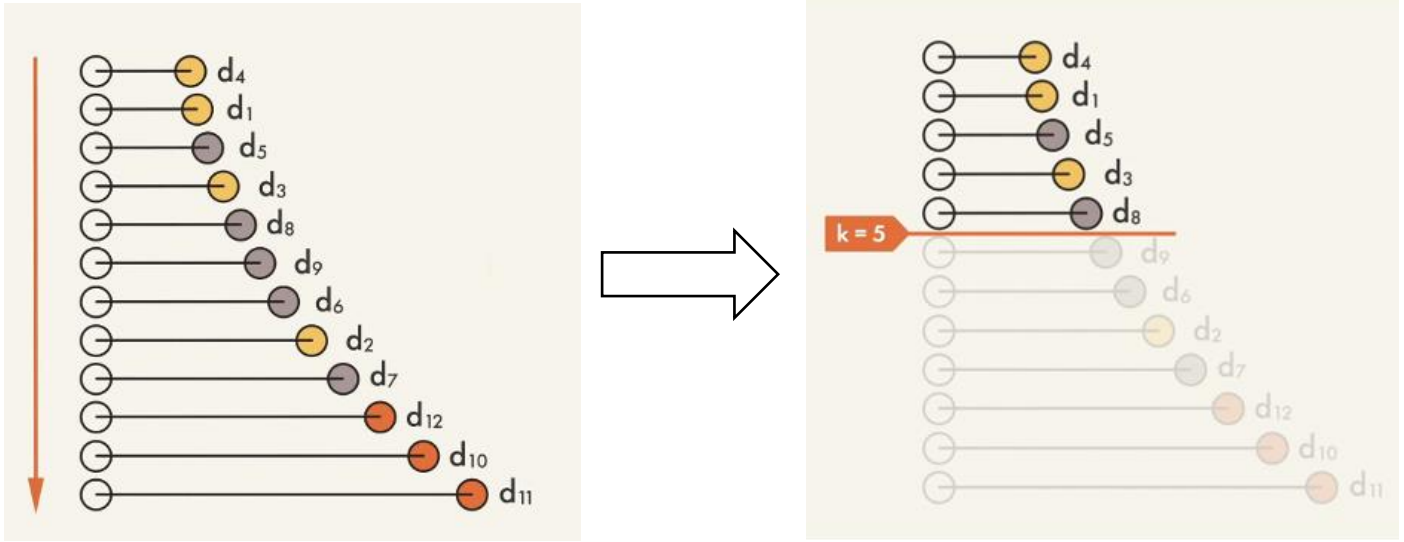

What could we have improved?

Our Project:

- 1 train data set
- 1 test data set

Improvement:

- ➢ Split of data sets
- 1 train data set
- 1 validation data set
- 1 test data set